

Advanced CNN


Grigory Sapunov  CTO / Intento

Использование CNN для других задач

- Трансформация изображений (перенос стиля, синтез изображений, ...)
- Обработка видео (spatio-temporal ...)
<https://arxiv.org/abs/1606.04698>
- Обработка пространственных данных
<http://www.microsoft.com/en-us/research/wp-content/uploads/2016/09/DeepST-SIGSPATIAL2016.pdf>
- Обработка звука (ASR, WaveNet, ...)
<http://benanne.github.io/2014/08/05/spotify-cnns.html>
<https://deepmind.com/blog/wavenet-generative-model-raw-audio/>
http://ronan.collobert.com/pub/matos/2015_cnnspeech_interspeech.pdf
<http://www.microsoft.com/en-us/research/publication/convolutional-neural-networks-for-speech-recognition-2/>
- Обработка текстов (классификация, перевод(!), ...)
<http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>
<https://arxiv.org/abs/1611.02344>


1D Convolution

1D




$$y[2] = x[0]w[4] + x[1]w[3] + x[2]w[2] + x[3]w[1] + x[4]w[0] = 1*1 + 2*1 + 3*2 + 4*1 + 5*1 = 18$$

2D



1D Convolution



<https://www.kaggle.com/shivamb/3d-convolutions-understanding-and-implementation>

1D Convolution Example: Classification


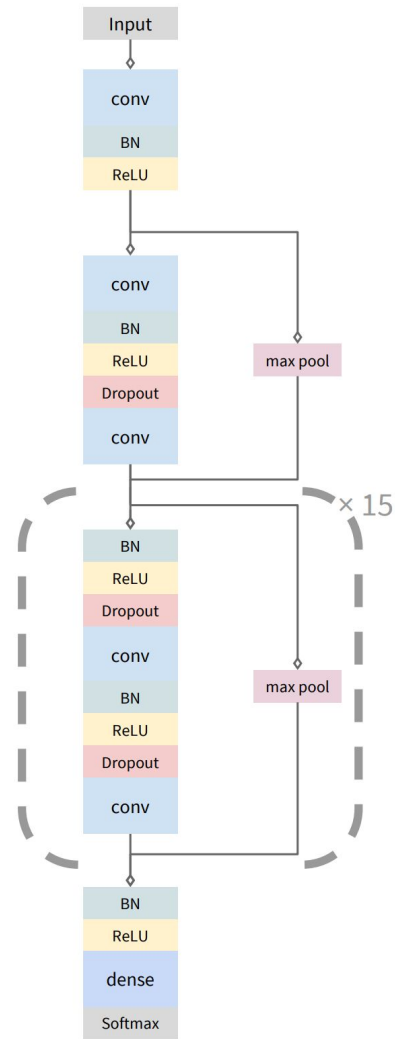


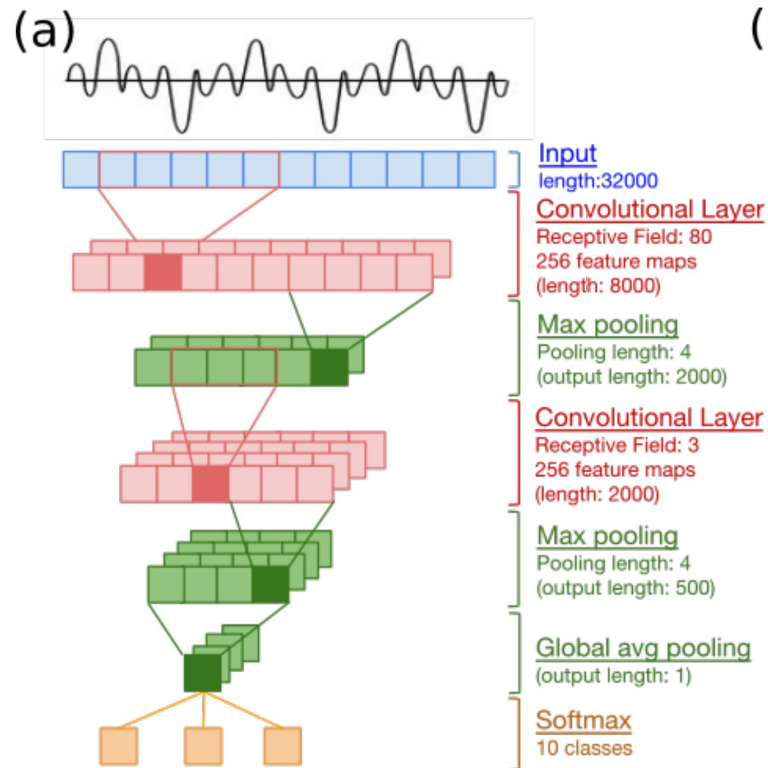
Figure 1. Our trained convolutional neural network correctly detecting the sinus rhythm (SINUS) and Atrial Fibrillation (AFIB) from this ECG recorded with a single-lead wearable heart monitor.



Cardiologist-Level Arrhythmia Detection with Convolutional Neural Networks, <https://arxiv.org/abs/1707.01836>

1D Convolution Example: Classification


M3 (0.2M)	M5 (0.5M)	M11 (1.8M)	M18 (3.7M)	M34-res (4M)
Input: 32000x1 time-domain waveform				
[80/4, 256]	[80/4, 128]	[80/4, 64]	[80/4, 64]	[80/4, 48]
Maxpool: 4x1 (output: 2000 × n)				
[3, 256]	[3, 128]	[3, 64] × 2	[3, 64] × 4	$\begin{bmatrix} 3, 48 \\ 3, 48 \end{bmatrix} \times 3$
Maxpool: 4x1 (output: 500 × n)				
	[3, 256]	[3, 128] × 2	[3, 128] × 4	$\begin{bmatrix} 3, 96 \\ 3, 96 \end{bmatrix} \times 4$
	Maxpool: 4x1 (output: 125 × n)			
	[3, 512]	[3, 256] × 3	[3, 256] × 4	$\begin{bmatrix} 3, 192 \\ 3, 192 \end{bmatrix} \times 6$
	Maxpool: 4x1 (output: 32 × n)			
		[3, 512] × 2	[3, 512] × 4	$\begin{bmatrix} 3, 384 \\ 3, 384 \end{bmatrix} \times 3$
Global average pooling (output: 1 × n)				
Softmax				



Very Deep Convolutional Neural Networks for Raw Waveforms,


<https://arxiv.org/abs/1610.00087>

3D Convolution




3D Convolution

3D conv.

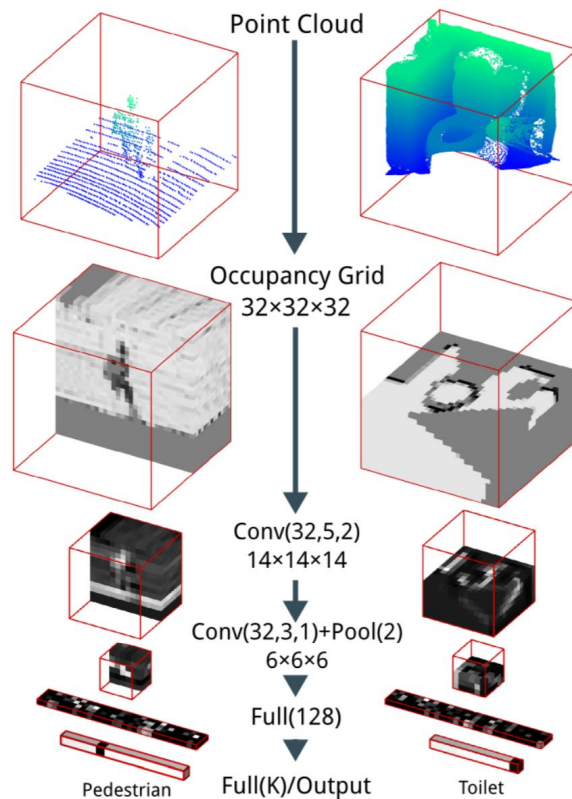


2D conv. with channels



<https://stackoverflow.com/questions/42883547/what-do-you-mean-by-1d-2d-and-3d-convolutions-in-cnn>


3D Convolution Example: Classification



VoxNet: A 3D Convolutional Neural Network for Real-Time Object Recognition

<https://www.ri.cmu.edu/publications/voxnet-a-3d-convolutional-neural-network-for-real-time-object-recognition/>

3D Convolution Example: Segmentation



3D Convolutional Neural Networks for Tumor Segmentation using Long-range 2D Context

<https://arxiv.org/abs/1807.08599>

Convolutions in Keras

- `keras.layers.Conv1D`

Input: 3D tensor with shape: (batch, steps, channels)

Output: 3D tensor with shape: (batch, new_steps, filters)

- `keras.layers.Conv2D`

2D convolution layer (e.g. spatial convolution over images).

Input: 4D tensor with shape: (batch, channels, rows, cols)

Output: 4D tensor with shape: (batch, filters, new_rows, new_cols)


- `keras.layers.Conv3D`

Input: 5D tensor with shape: (batch, channels, conv_dim1, conv_dim2, conv_dim3)

Output: 5D tensor with shape: (batch, filters, new_conv_dim1, new_conv_dim2, new_conv_dim3)

<https://keras.io/layers/convolutional/>

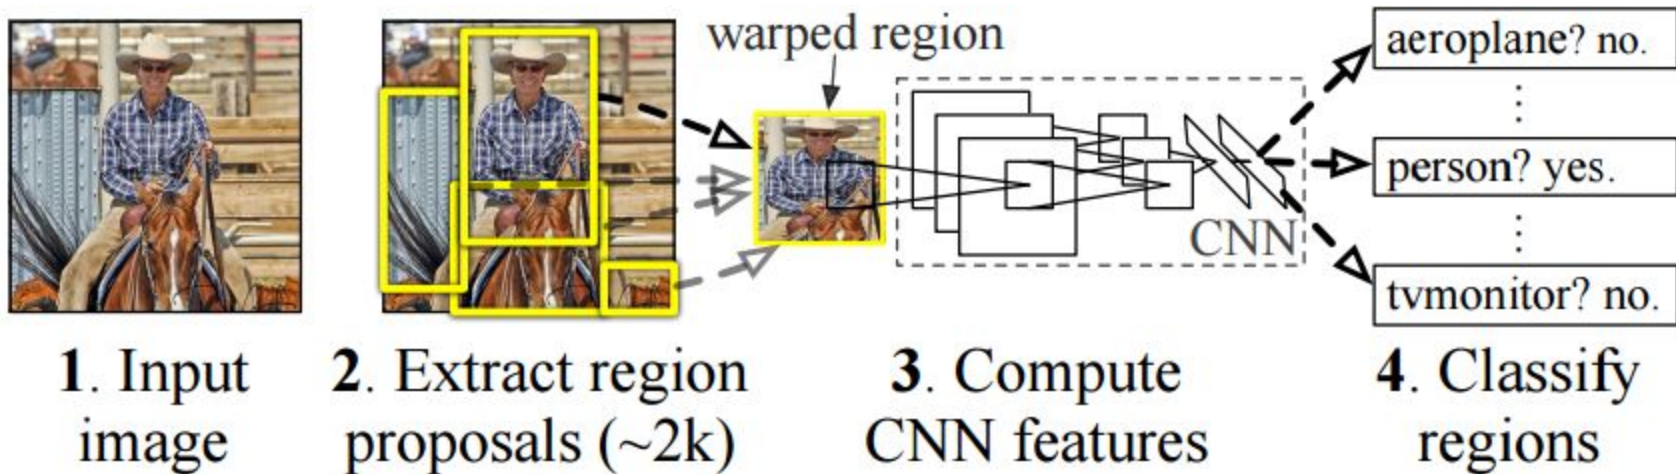
Dilated/Atrous Convolutions



Детекция:

R-CNN, Fast R-CNN, Faster R-CNN

R-CNN: Region-based Convolutional Network




<http://nbviewer.jupyter.org/github/BVLC/caffe/blob/master/examples/detection.ipynb>

<https://github.com/rbgirshick/rcnn>

https://people.eecs.berkeley.edu/~rbg/papers/pami/rcnn_pami.pdf


Fast R-CNN



<http://tutorial.caffe.berkeleyvision.org/caffe-cvpr15-detection.pdf>

Faster R-CNN

Key Idea: Region Proposal Net (RPN) layer




<https://arxiv.org/abs/1506.01497>

https://github.com/ShaoqingRen/faster_rcnn

<http://www.slideshare.net/ssuser416c44/faster-rcnn>

YOLO: Real-Time Object Detection




<https://pjreddie.com/darknet/yolo/>

Fully-convolutional networks (FCN)

Fully-convolutional networks (FCN)

Обычная свёрточная сеть, но без MLP сверху (нет полносвязных слоёв).

Позволяет **работать с изображениями произвольного размера** и выдавать на выходе тепловую карту классификации.



Fully-convolutional networks (FCN)

Можно преобразовать полносвязный слой на свёрточный

```
layer {  
  name: "fc6"  
  type: "InnerProduct"  
  bottom: "pool5"  
  top: "fc6"  
  inner_product_param {  
    num_output: 4096  
  }  
}
```

```
layer {  
  name: "conv6"  
  type: "Convolution"  
  bottom: "pool5"  
  top: "conv6"  
  convolution_param {  
    num_output: 4096  
    kernel_size: 6  
  }  
}
```

Table 1: Left: fc6 definition, Right: equivalent conv6 definition with a kernel size of 6 because the input to fc6 is a 6×6 image patch.

Ресурсы

- https://people.eecs.berkeley.edu/~jonlong/long_shelhamer_fcn.pdf
- <https://github.com/shelhamer/fcn.berkeleyvision.org>
- <https://github.com/developmentseed/caffe-fcn>


- <https://www.quora.com/How-is-Fully-Convolutional-Network-FCN-different-from-the-original-Convolutional-Neural-Network-CNN>
- <https://www.quora.com/What-are-the-advantages-of-Fully-Convolutional-Networks-over-CNNs>
- <https://www.quora.com/How-does-the-conversion-of-last-layers-of-CNN-from-fully-connected-to-fully-convolutional-allow-it-to-process-images-of-different-size>

Deconvolution networks

Deconvolution networks

Правильнее называть это Transposed convolution, а не Deconvolution (это слово уже занято в цифровой обработке сигналов для обратной операции).

По сути, реализован обучаемый upsampling.



Deconvolution networks


Специальный тип слоя: Deconvolution.

Параметр stride задаёт степень увеличения.


```
layer {  
  name: "upscore"  
  type: "Deconvolution"  
  bottom: "score"  
  top: "upscore"  
  convolution_param {  
    num_output: 12 # set this to number of classes  
    kernel_size: 63  
    stride: 32  
  }  
}
```


Как работает Transposed convolution

- 1) Делаем padding нулями (unpooling)
- 2) Применяем convolution



Генерация изображений




Ресурсы

- <https://devblogs.nvidia.com/parallelforall/image-segmentation-using-digits-5>
- https://github.com/vdumoulin/conv_arithmetic
- A guide to convolution arithmetic for deep learning
<https://arxiv.org/abs/1603.07285>
- <https://towardsdatascience.com/types-of-convolutions-in-deep-learning-717013397f4d>
- <https://towardsdatascience.com/transpose-convolution-77818e55a123>
- <https://www.quora.com/How-does-a-deconvolutional-neural-network-work>
- <http://stackoverflow.com/questions/35049197/how-does-the-unpooling-and-deconvolution-work-in-deconvnet>
- Кейс с убиением очков
<https://blog.insightdatascience.com/isee-removing-eyeglasses-from-faces-using-deep-learning-d4e7d935376f>

Generative Adversarial Networks (GANs)

Generative Adversarial Networks (GANs)



Generative Adversarial Networks (GANs)




Figure 1: DCGAN generator used for LSUN scene modeling. A 100 dimensional uniform distribution Z is projected to a small spatial extent convolutional representation with many feature maps. A series of four fractionally-strided convolutions (in some recent papers, these are wrongly called deconvolutions) then convert this high level representation into a 64×64 pixel image. Notably, no fully connected or pooling layers are used.

Generative Adversarial Networks (GANs)



Figure 2: Generated bedrooms after one training pass through the dataset. Theoretically, the model could learn to memorize training examples, but this is experimentally unlikely as we train with a small learning rate and minibatch SGD. We are aware of no prior empirical evidence demonstrating memorization with SGD and a small learning rate.

<https://arxiv.org/pdf/1511.06434v2.pdf>

Generative Adversarial Networks (GANs)



volcano

<http://www.evolvingai.org/ppgn>

Example: StackGAN

This small blue bird has a short pointy beak and brown on its wings



This bird is completely red with black wings and pointy beak



A small sized bird that has a cream belly and a short pointed bill



A small bird with a black head and wings and features grey wings



StackGAN: Text to Photo-realistic Image Synthesis with
Stacked Generative Adversarial Networks, <https://arxiv.org/abs/1612.03242>

Example: Progressive Growing of GANs




Figure 1: Our training starts with both the generator (G) and discriminator (D) having a low spatial resolution of 4×4 pixels. As the training advances, we incrementally add layers to G and D, thus increasing the spatial resolution of the generated images. All existing layers remain trainable throughout the process. Here $N \times N$ refers to convolutional layers operating on $N \times N$ spatial resolution. This allows stable synthesis in high resolutions and also speeds up training considerably. On the right we show six example images generated using progressive growing at 1024×1024 .

Progressive Growing of GANs for Improved Quality, Stability, and Variation,

<https://arxiv.org/abs/1710.10196>

<https://www.youtube.com/watch?v=XOxxPcy5Gr4>

Ресурсы

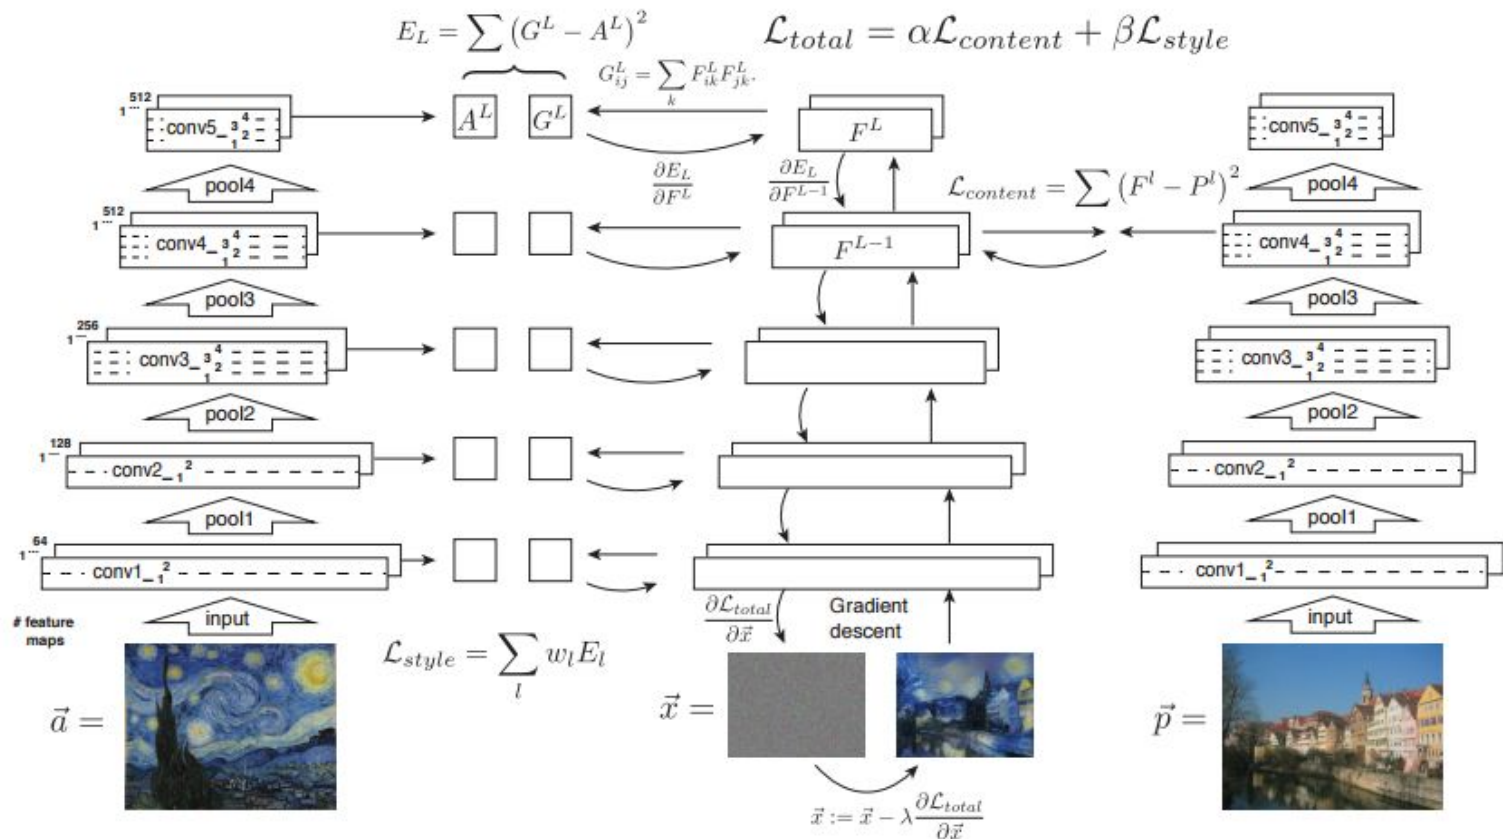
- <https://openai.com/blog/generative-models/>
- <https://www.quora.com/What-are-Generative-Adversarial-Networks>
- Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks
<https://arxiv.org/abs/1511.06434>
- <http://blog.aylien.com/introduction-generative-adversarial-networks-code-tens-orflow/>
- Plug & Play Generative Networks
<http://www.evolvingai.org/ppgn>

Кейс: Перенос стиля

Неклассические задачи: перенос стиля

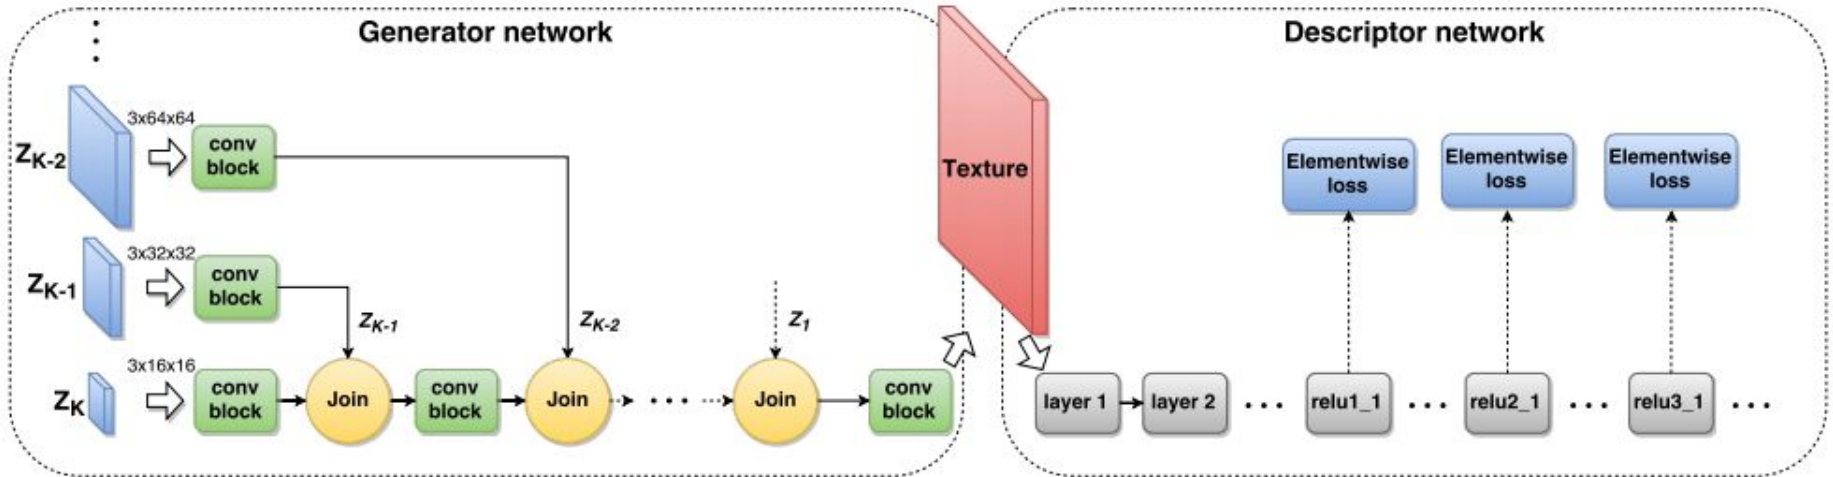


Перенос стиля: оригинальный алгоритм



Перенос стиля: быстрый алгоритм

Texture Networks



<https://arxiv.org/abs/1603.03417>

<https://research.googleblog.com/2016/10/supercharging-style-transfer.html>

Реализации

Классический алгоритм:

- <https://github.com/jcjohnson/neural-style> (Torch)
- <https://github.com/fzliu/style-transfer> (Caffe)
- <https://github.com/titu1994/Neural-Style-Transfer> (Keras)

Быстрые алгоритмы:

- <https://github.com/jcjohnson/fast-neural-style> (Torch)
- <https://github.com/yusuketomoto/chainer-fast-neuralstyle> (Python/Chainer)
- <https://github.com/lengstrom/fast-style-transfer> (Python/TF)
- https://github.com/DmitryUlyanov/texture_nets (Torch)

Doodle:

- <https://github.com/alexjc/neural-doodle>